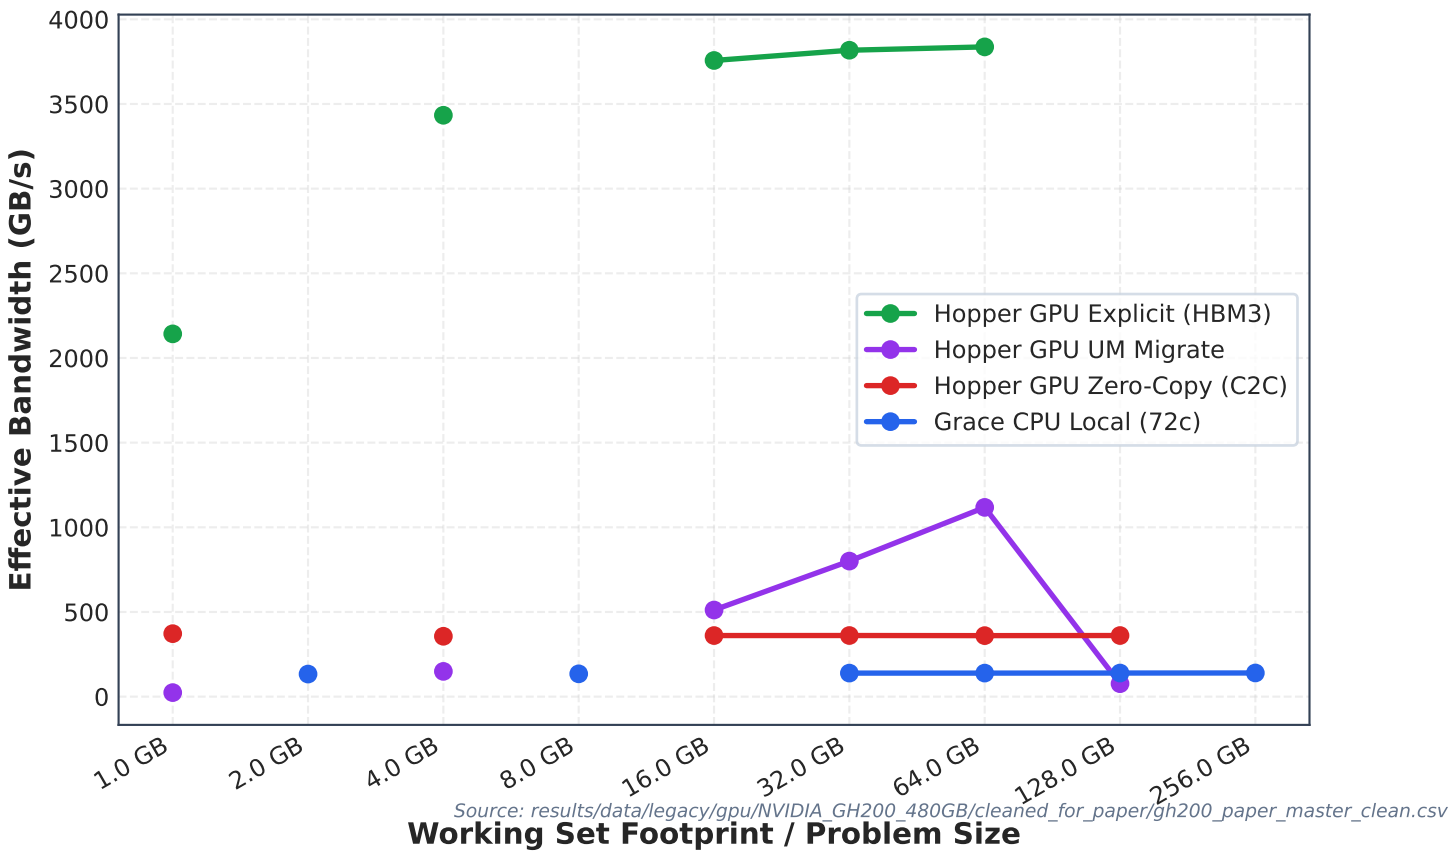
LLAMA3 KV CACHE OFFLOAD — Throughput vs Footprint by Residency Mode
Effective Bandwidth (GB/s)
4000
3500
3000
2500
2000
1500
1000
500
0
Hopper GPU Explicit (HBM3)
Hopper GPU UM Migrate
Hopper GPU Zero-Copy (C2C)
Grace CPU Local (72c)
1.0 GB
2.0 GB
4.0 GB
8.0 GB
16.0 GB
32.0 GB
64.0 GB
128.0 GB
256.0 GB
Source: results/data/legacy/gpu/NVIDIA_GH200_480GB/cleaned_for_paper/gh200_paper_master_clean.csv
Working Set Footprint / Problem Size