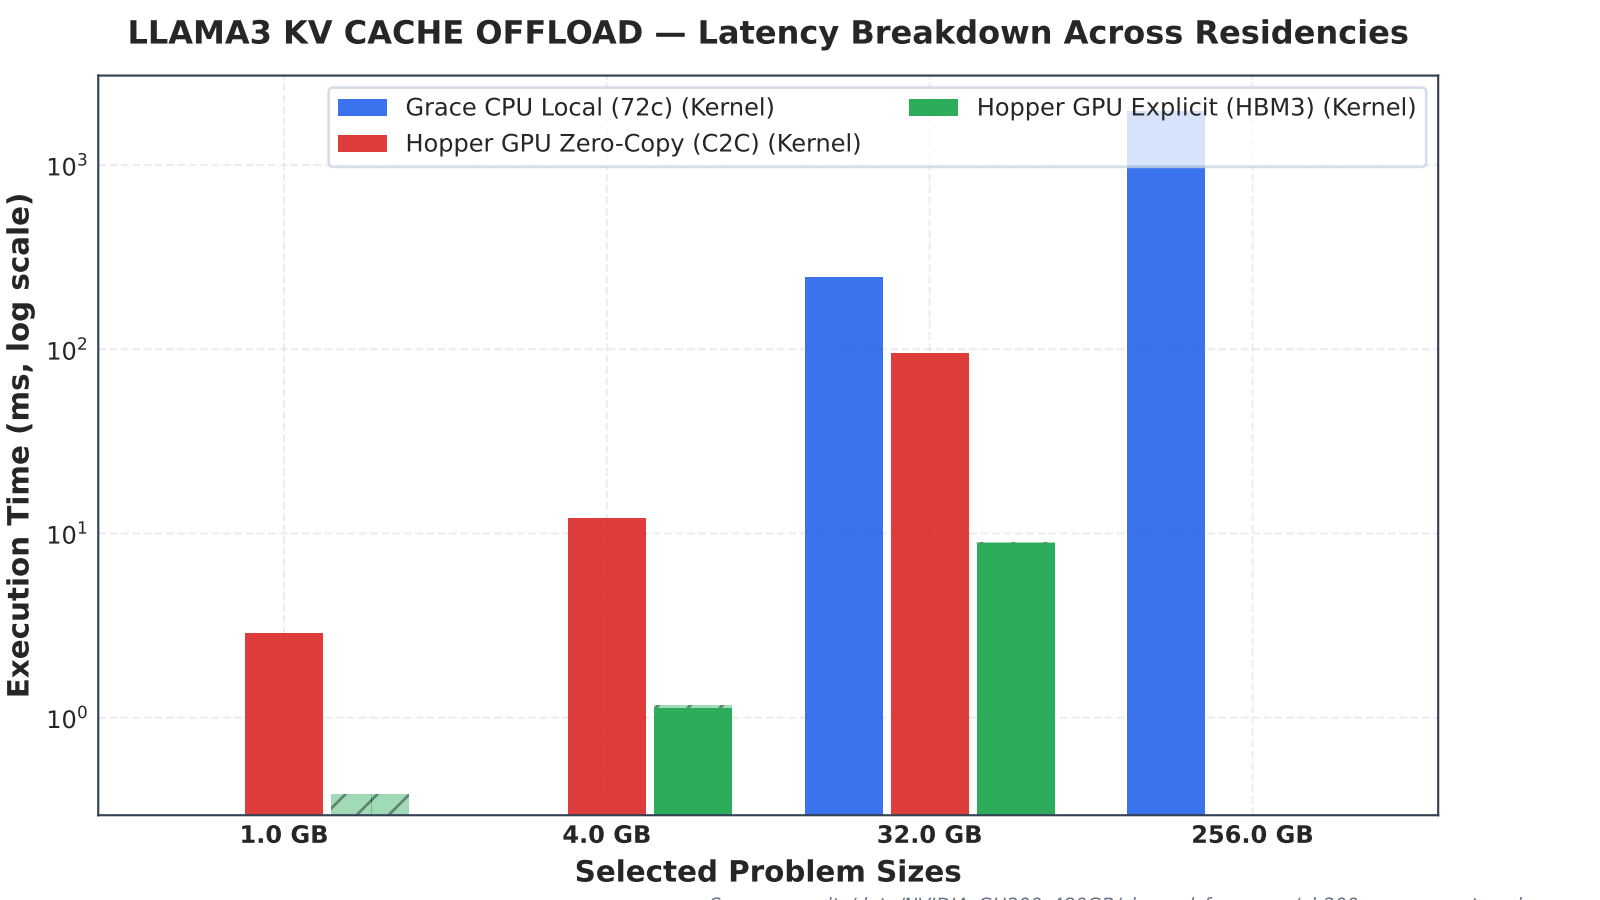
LLAMA3 KV CACHE OFFLOAD — Latency Breakdown Across Residencies
Grace CPU Local (72c) (Kernel)
Hopper GPU Zero-Copy (C2C) (Kernel)
Hopper GPU Explicit (HBM3) (Kernel)
Execution Time (ms, log scale)
10^3
10^2
10^1
10^0
1.0 GB
4.0 GB
32.0 GB
256.0 GB
Selected Problem Sizes
Source: results/data/NVIDIA_GH200_480GB/cleaned_for_paper/gh200_paper_master_clean.csv