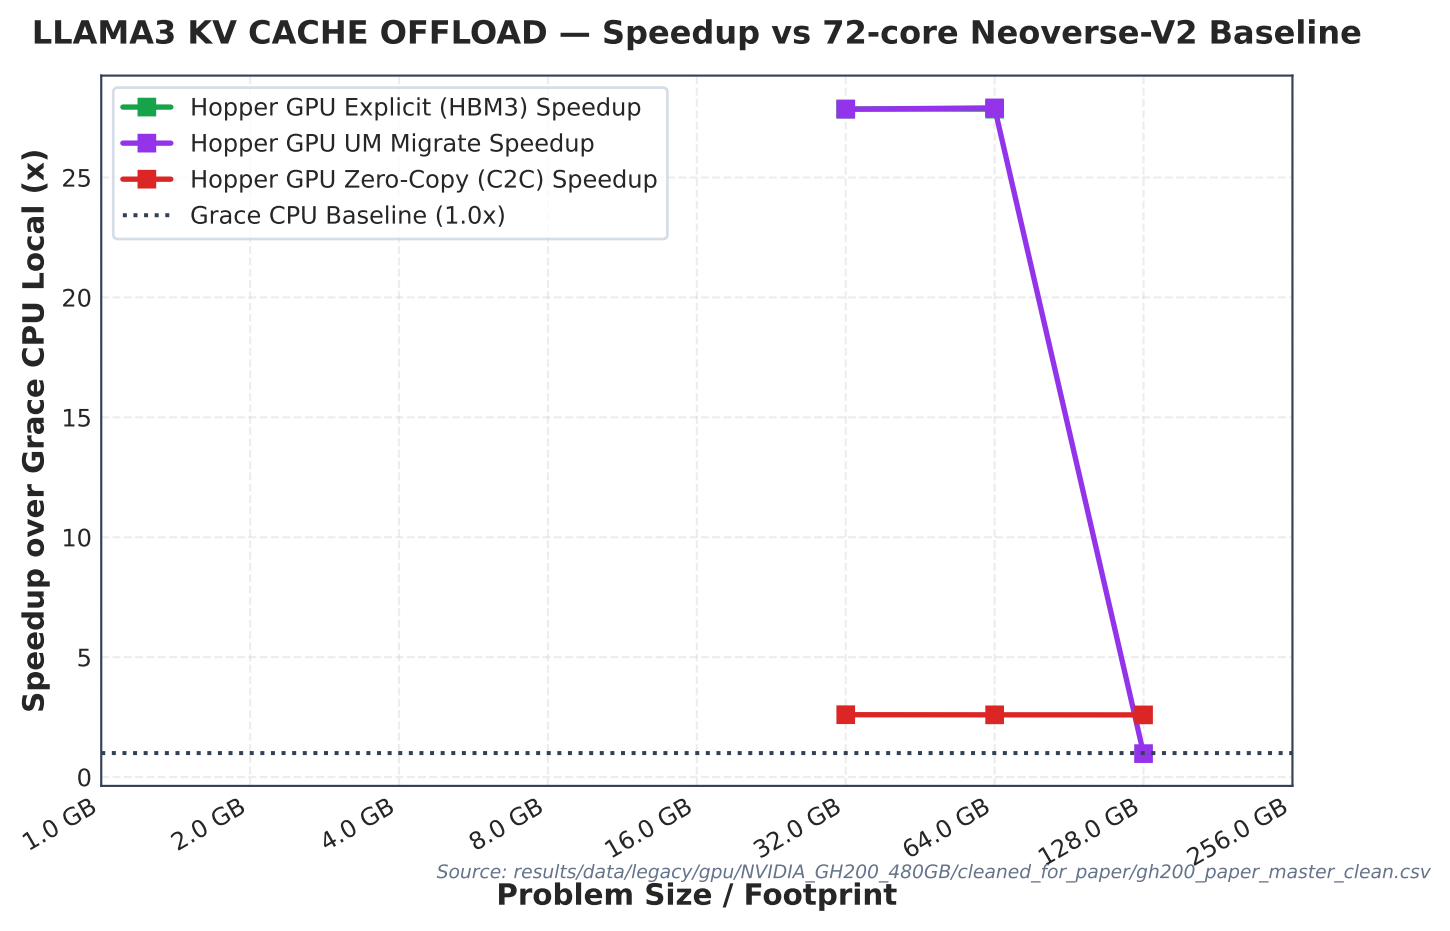
LLAMA3 KV CACHE OFFLOAD — Speedup vs 72-core Neoverse-V2 Baseline
Hopper GPU Explicit (HBM3) Speedup
Hopper GPU UM Migrate Speedup
Hopper GPU Zero-Copy (C2C) Speedup
Grace CPU Baseline (1.0x)
Speedup over Grace CPU Local (x)
0
5
10
15
20
25
1.0 GB
2.0 GB
4.0 GB
8.0 GB
16.0 GB
32.0 GB
64.0 GB
128.0 GB
256.0 GB
Problem Size / Footprint
Source: results/data/legacy/gpu/NVIDIA_GH200_480GB/cleaned_for_paper/gh200_paper_master_clean.csv