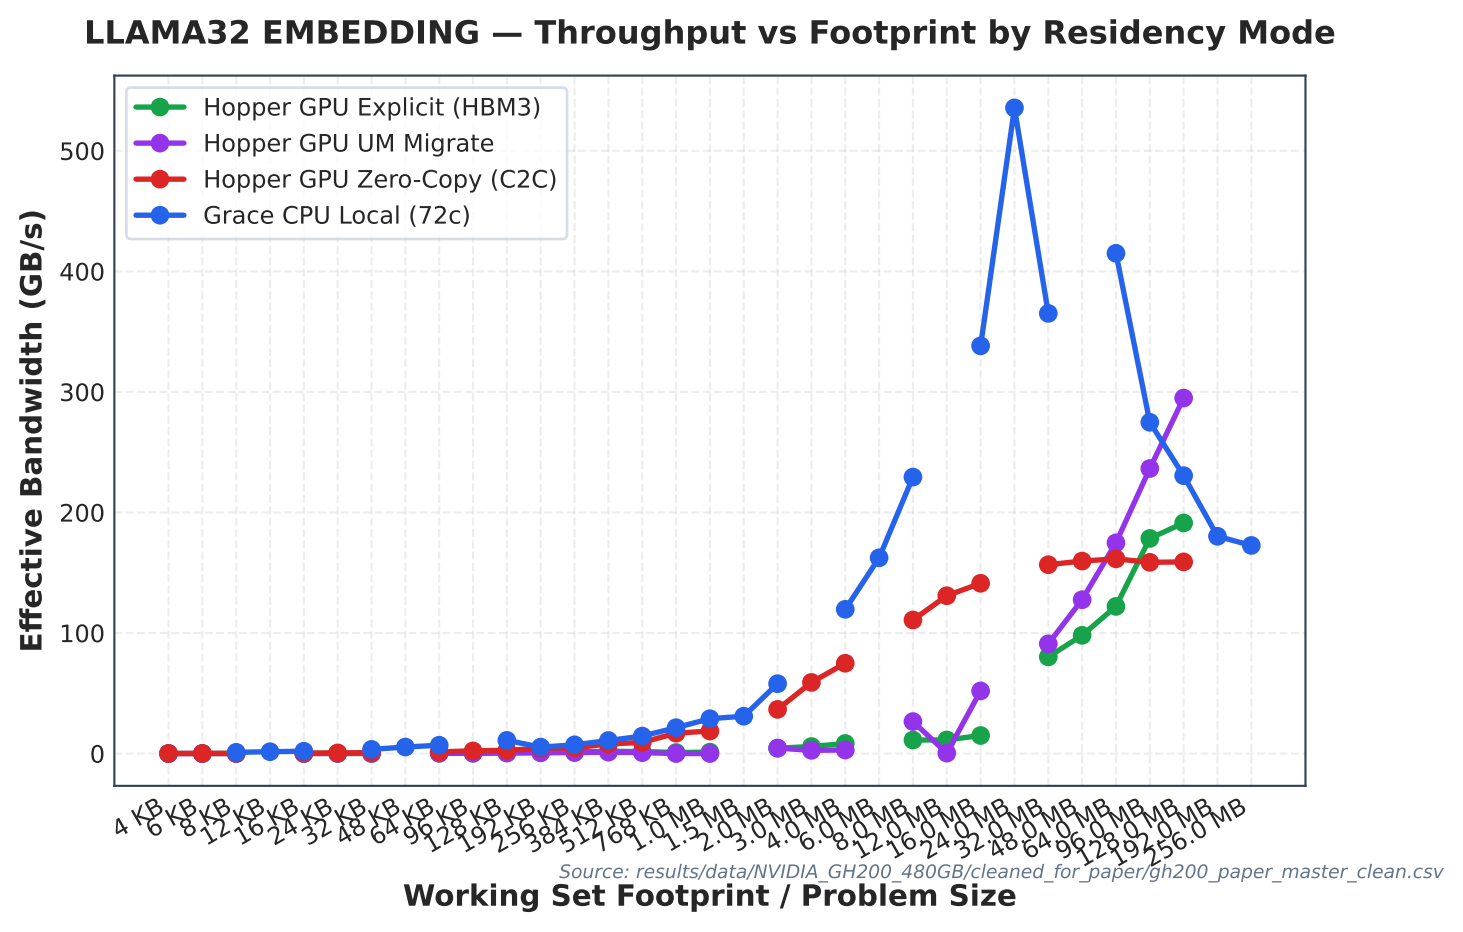
LLAMA32 EMBEDDING — Throughput vs Footprint by Residency Mode
Hopper GPU Explicit (HBM3)
Hopper GPU UM Migrate
Hopper GPU Zero-Copy (C2C)
Grace CPU Local (72c)
Effective Bandwidth (GB/s)
0
100
200
300
400
500
4 KB
6 KB
8 KB
12 KB
16 KB
24 KB
32 KB
48 KB
64 KB
96 KB
128 KB
192 KB
256 KB
384 KB
512 KB
768 KB
1.0 MB
1.5 MB
2.0 MB
3.0 MB
4.0 MB
6.0 MB
8.0 MB
12.0 MB
16.0 MB
24.0 MB
32.0 MB
48.0 MB
64.0 MB
96.0 MB
128.0 MB
192.0 MB
256.0 MB
Source: results/data/NVIDIA_GH200_480GB/cleaned_for_paper/gh200_paper_master_clean.csv
Working Set Footprint / Problem Size