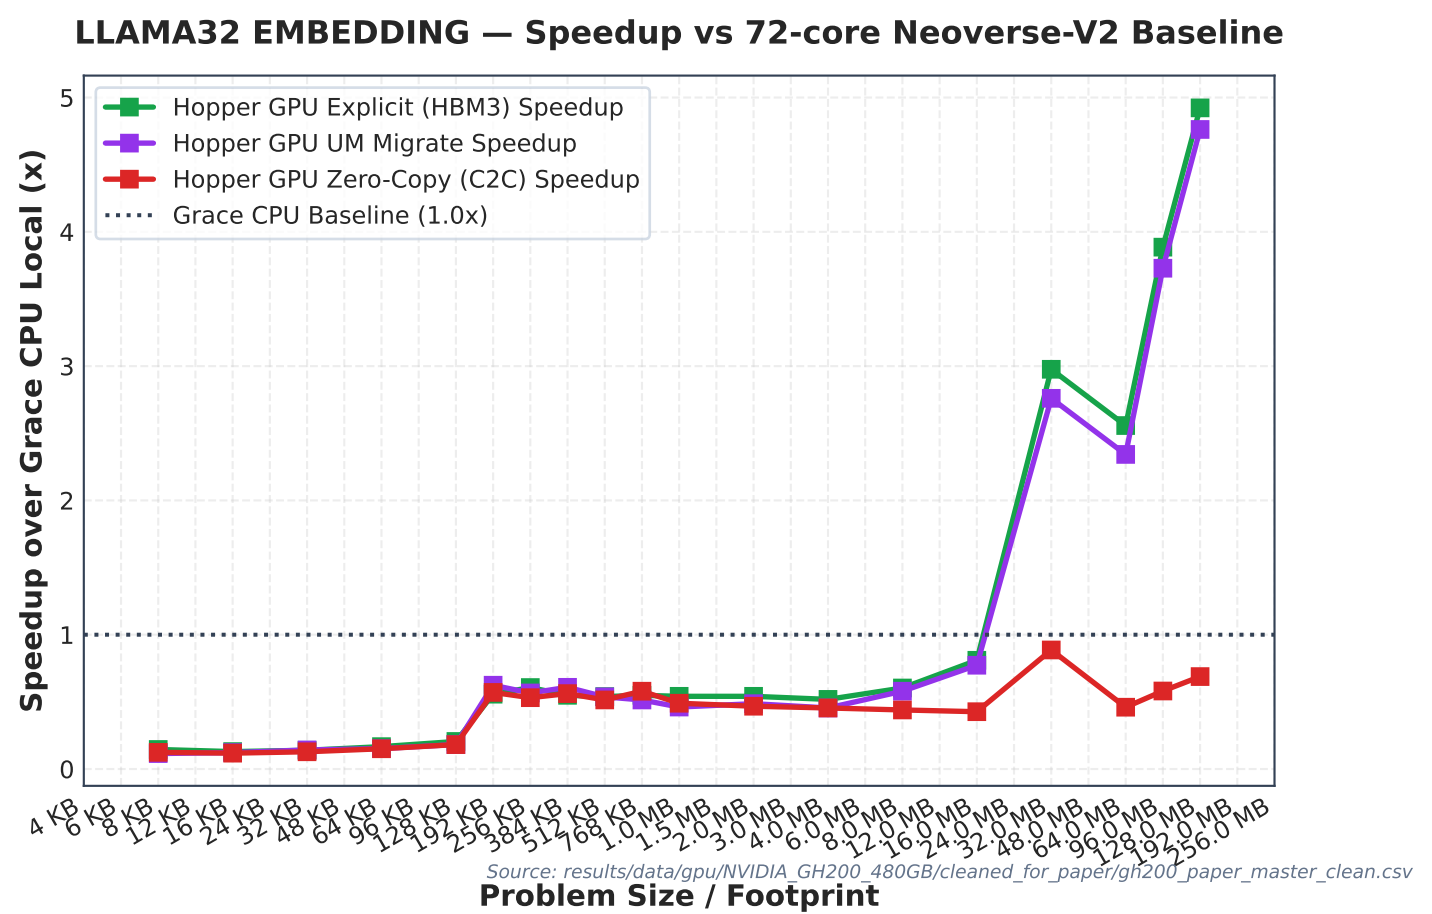
LLAMA32 EMBEDDING — Speedup vs 72-core Neoverse-V2 Baseline
Hopper GPU Explicit (HBM3) Speedup
Hopper GPU UM Migrate Speedup
Hopper GPU Zero-Copy (C2C) Speedup
Grace CPU Baseline (1.0x)
Speedup over Grace CPU Local (x)
0
1
2
3
4
5
4 KB
6 KB
8 KB
12 KB
16 KB
24 KB
32 KB
48 KB
64 KB
96 KB
128 KB
192 KB
256 KB
384 KB
512 KB
768 KB
1.0 MB
1.5 MB
2.0 MB
3.0 MB
4.0 MB
6.0 MB
8.0 MB
12.0 MB
16.0 MB
24.0 MB
32.0 MB
48.0 MB
64.0 MB
96.0 MB
128.0 MB
192.0 MB
256.0 MB
Source: results/data/gpu/NVIDIA_GH200_480GB/cleaned_for_paper/gh200_paper_master_clean.csv
Problem Size / Footprint