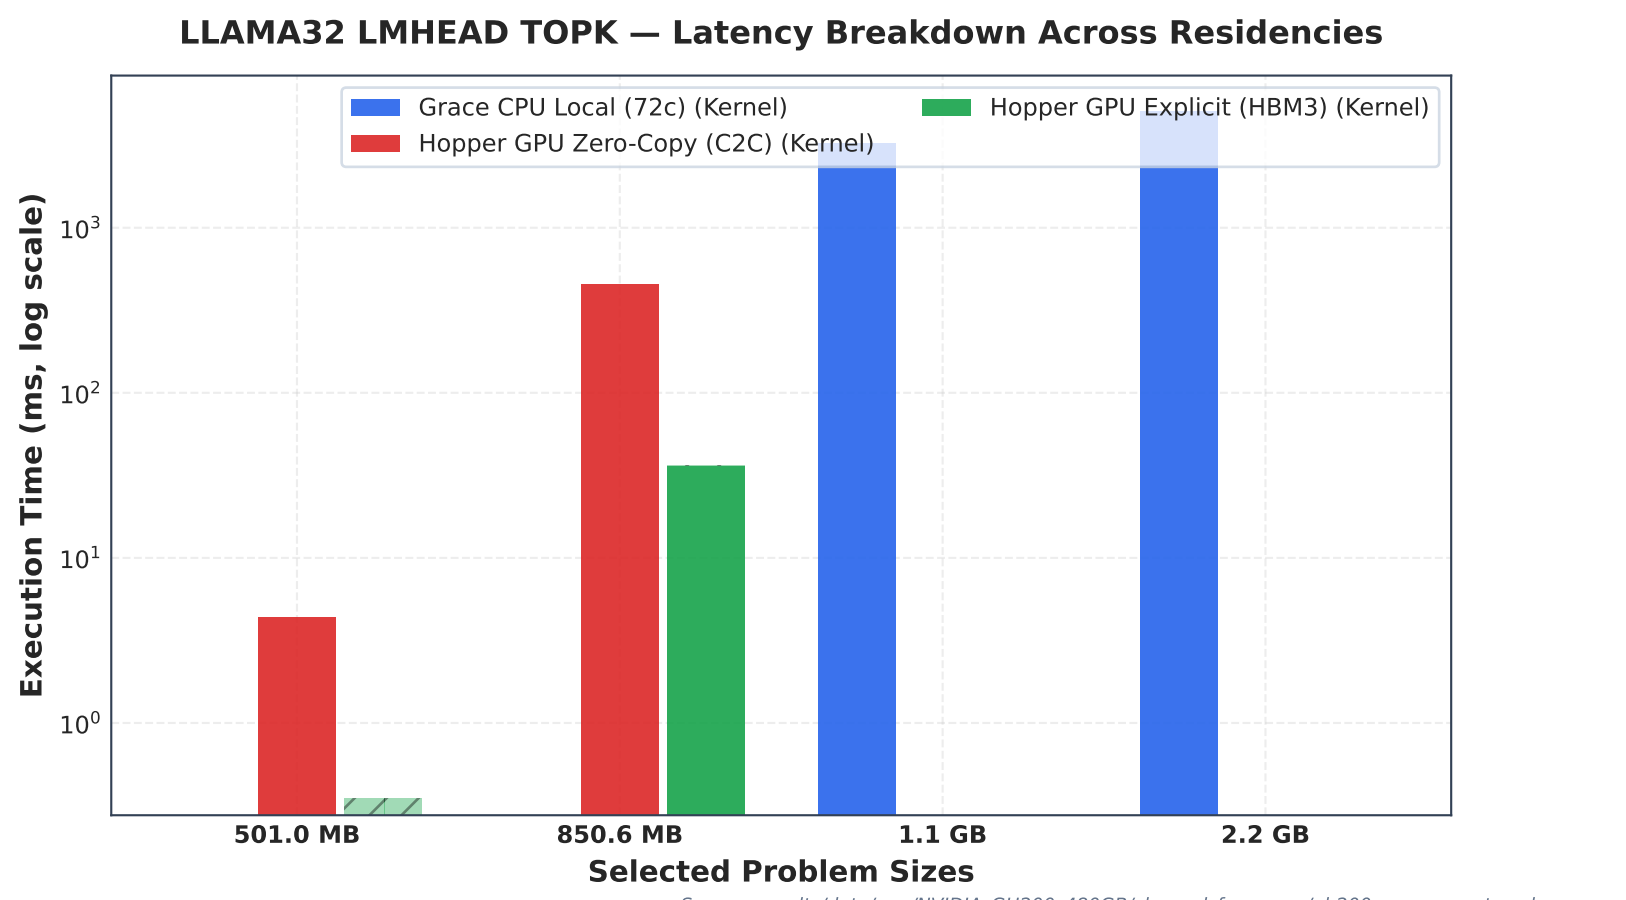
LLAMA32 LMHEAD TOPK — Latency Breakdown Across Residencies
Grace CPU Local (72c) (Kernel)
Hopper GPU Zero-Copy (C2C) (Kernel)
Hopper GPU Explicit (HBM3) (Kernel)
Execution Time (ms, log scale)
10^3
10^2
10^1
10^0
501.0 MB
850.6 MB
1.1 GB
2.2 GB
Selected Problem Sizes
Source: results/data/gpu/NVIDIA_GH200_480GB/cleaned_for_paper/gh200_paper_master_clean.csv