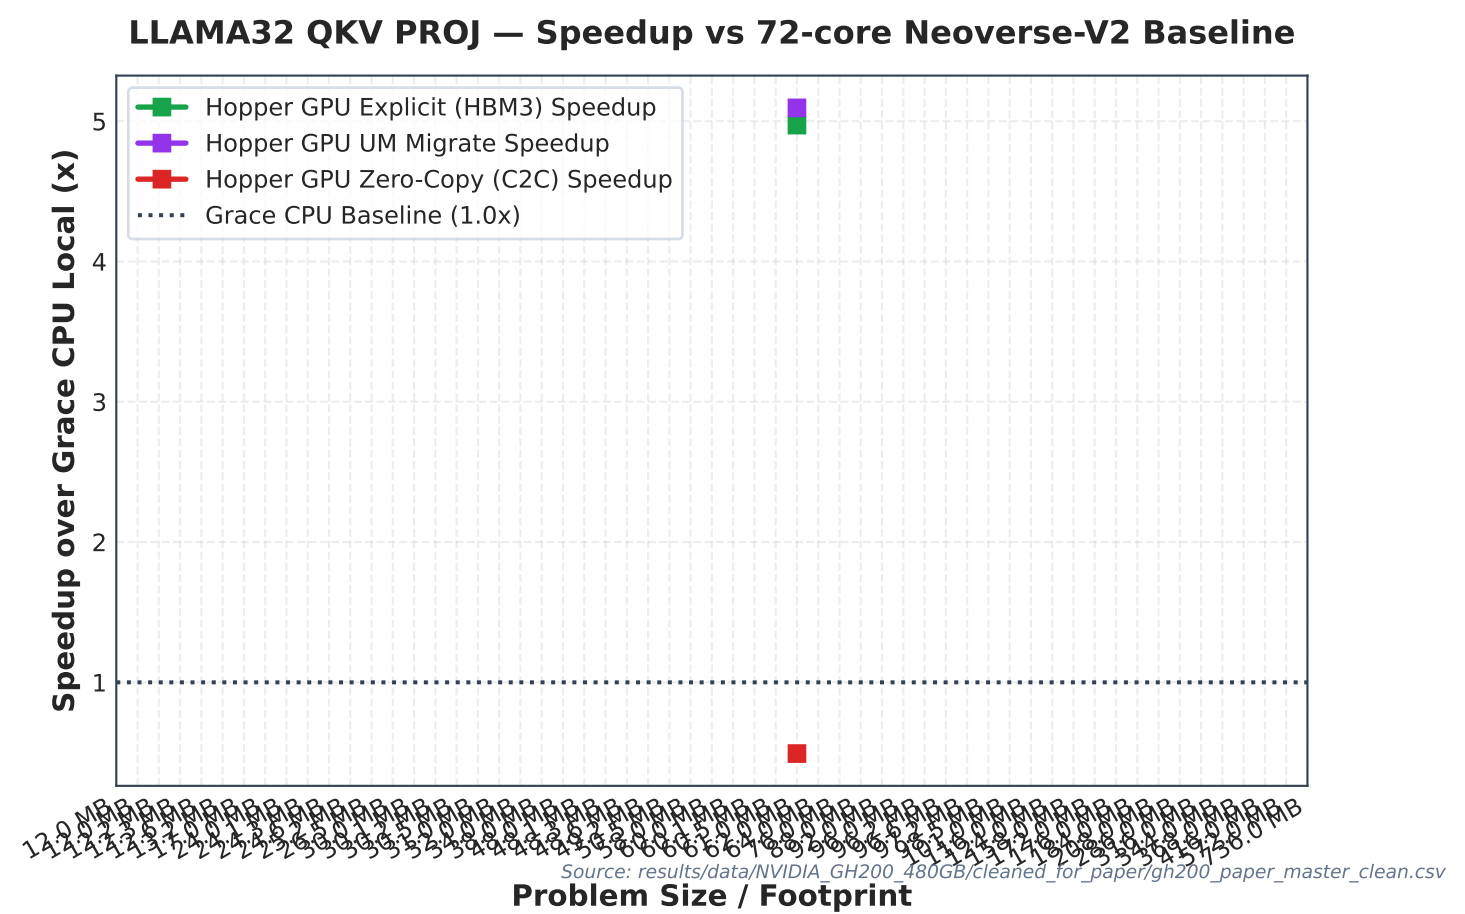
LLAMA32 QKV PROJ — Speedup vs 72-core Neoverse-V2 Baseline
Hopper GPU Explicit (HBM3) Speedup
Hopper GPU UM Migrate Speedup
Hopper GPU Zero-Copy (C2C) Speedup
Grace CPU Baseline (1.0x)
Speedup over Grace CPU Local (x)
1
2
3
4
5
Problem Size / Footprint
Source: results/data/NVIDIA_GH200_480GB/cleaned_for_paper/gh200_paper_master_clean.csv