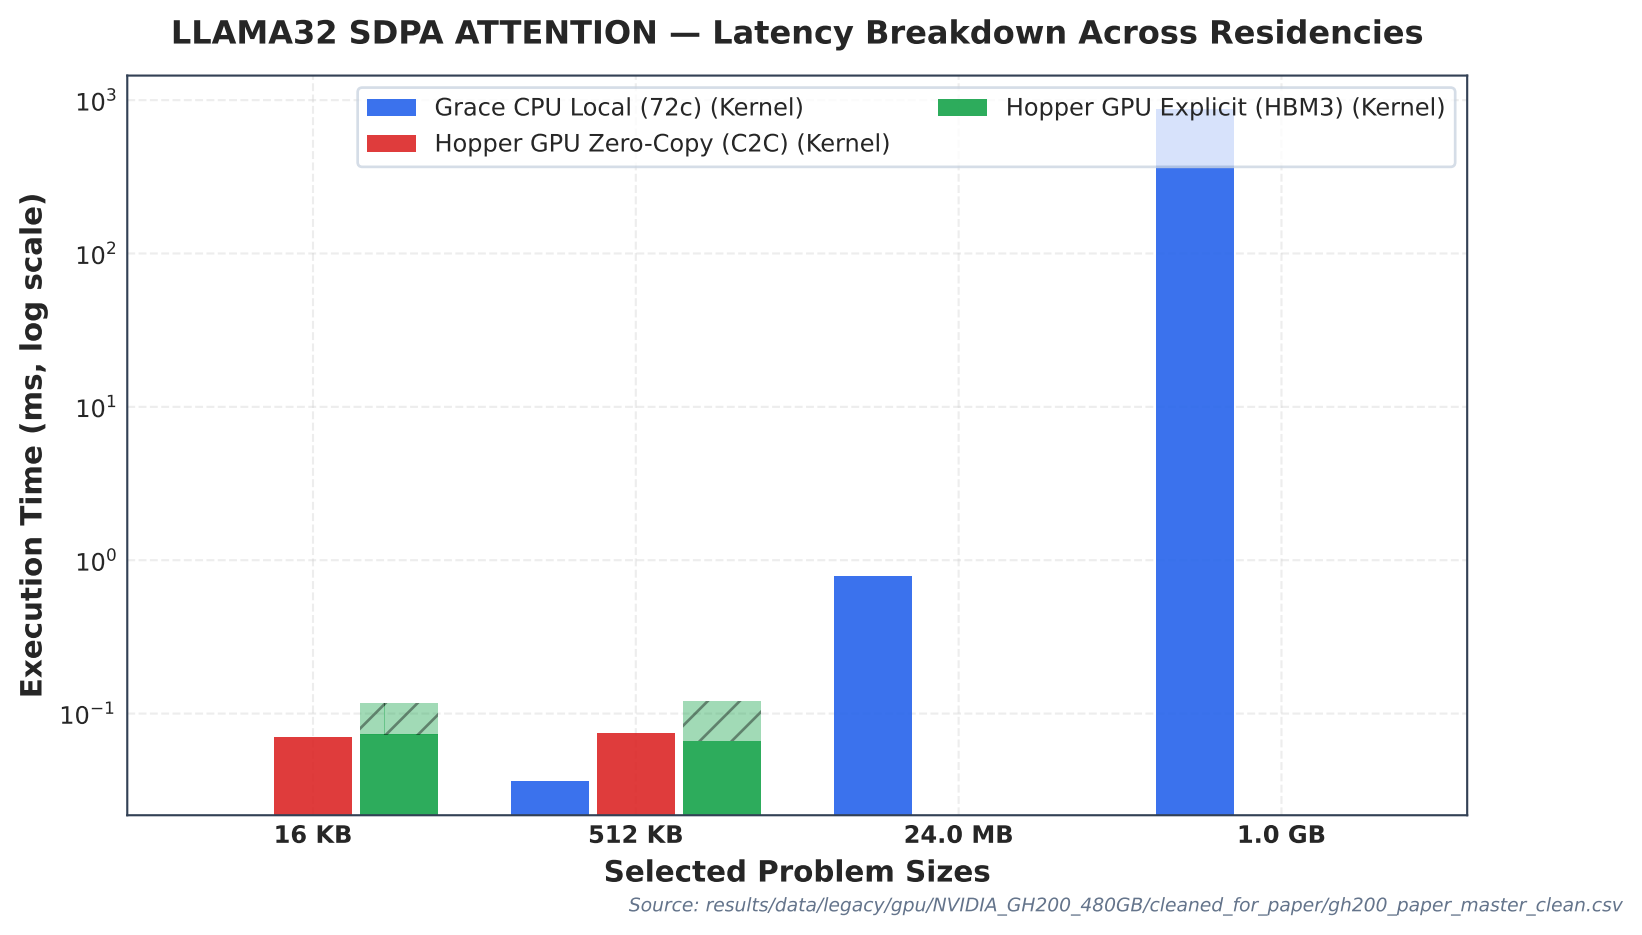
LLAMA32 SDPA ATTENTION — Latency Breakdown Across Residencies
Grace CPU Local (72c) (Kernel)
Hopper GPU Explicit (HBM3) (Kernel)
Hopper GPU Zero-Copy (C2C) (Kernel)
Execution Time (ms, log scale)
10^3
10^2
10^1
10^0
10^{-1}
16 KB
512 KB
24.0 MB
1.0 GB
Selected Problem Sizes
Source: results/data/legacy/gpu/NVIDIA_GH200_480GB/cleaned_for_paper/gh200_paper_master_clean.csv