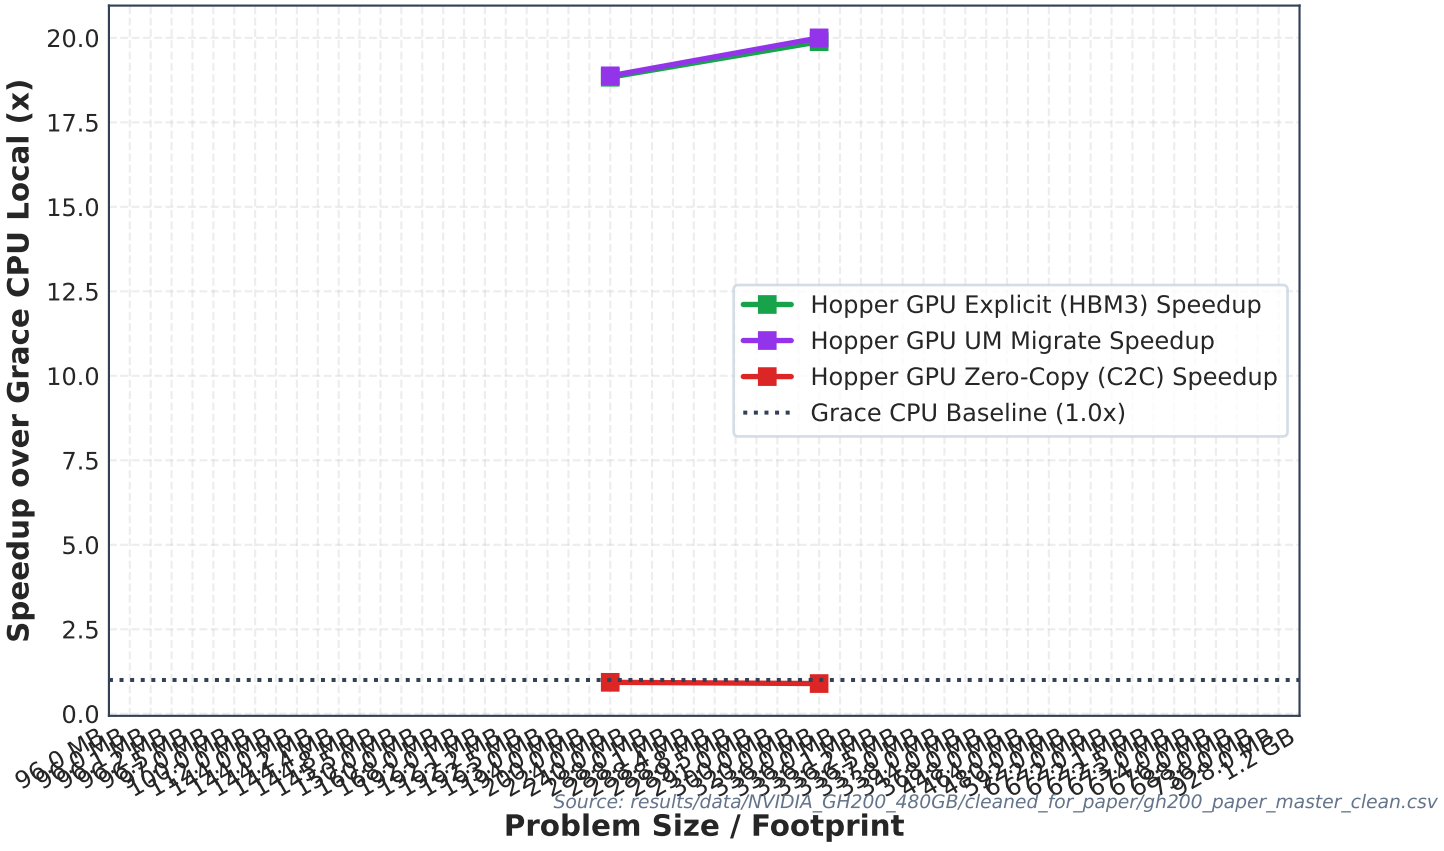
LLAMA32 SWIGLU FFN — Speedup vs 72-core Neoverse-V2 Baseline
Speedup over Grace CPU Local (x)
20.0
17.5
15.0
12.5
10.0
7.5
5.0
2.5
0.0
Hopper GPU Explicit (HBM3) Speedup
Hopper GPU UM Migrate Speedup
Hopper GPU Zero-Copy (C2C) Speedup
Grace CPU Baseline (1.0x)
960.0 MB
1.2 GB
Problem Size / Footprint
Source: results/data/NVIDIA_GH200_480GB/cleaned_for_paper/gh200_paper_master_clean.csv