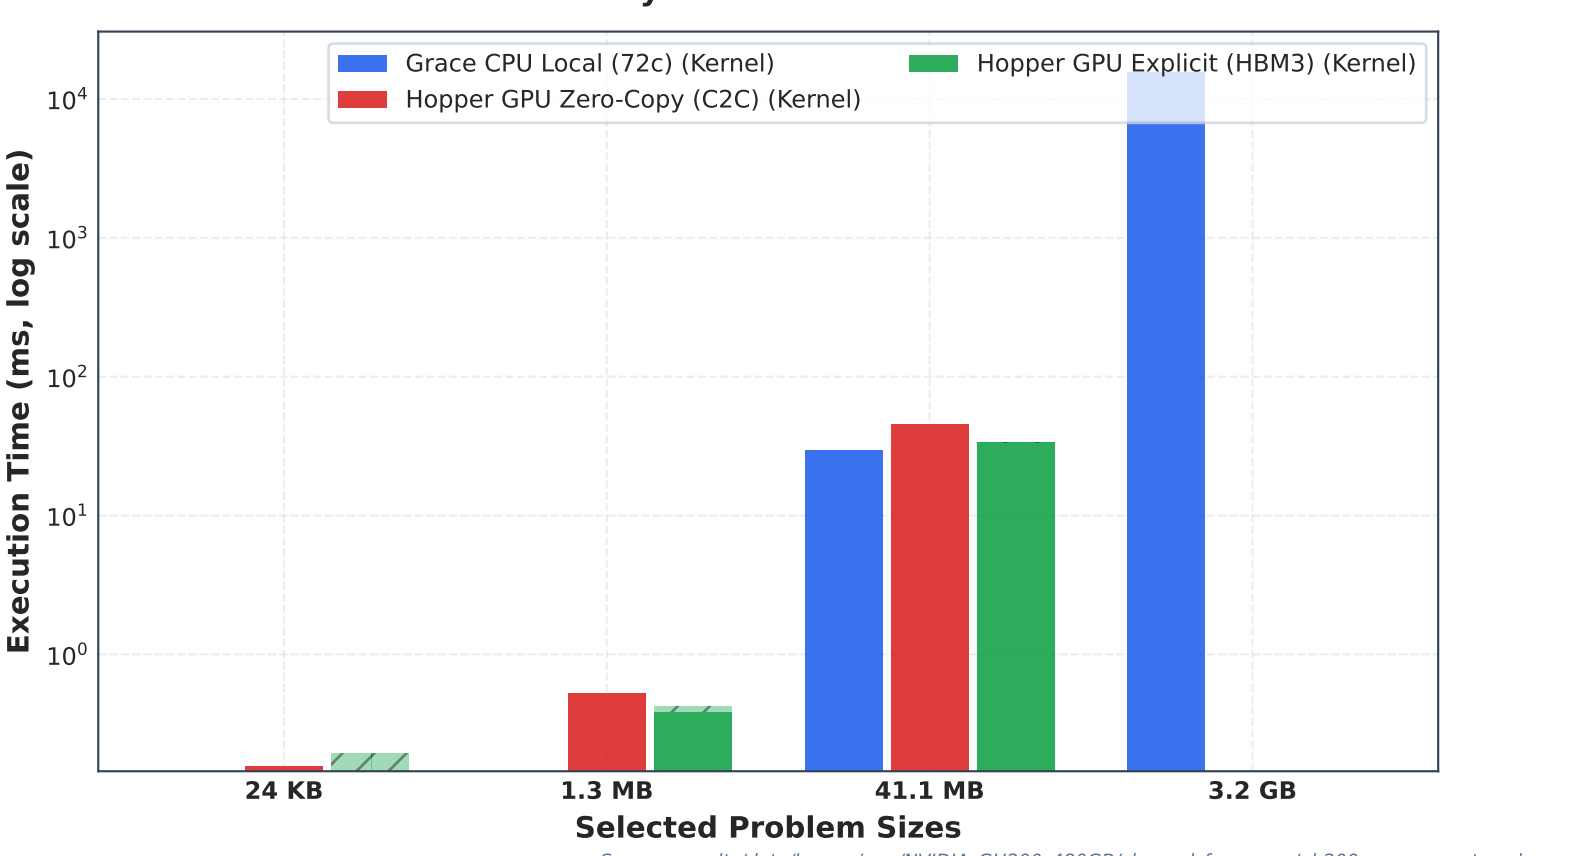
SPGEMM — Latency Breakdown Across Residencies
Grace CPU Local (72c) (Kernel)
Hopper GPU Zero-Copy (C2C) (Kernel)
Hopper GPU Explicit (HBM3) (Kernel)
Execution Time (ms, log scale)
$10^4$
$10^3$
$10^2$
$10^1$
$10^0$
24 KB
1.3 MB
41.1 MB
3.2 GB
Selected Problem Sizes
Source: results/data/legacy/gpu/NVIDIA_GH200_480GB/cleaned_for_paper/gh200_paper_master_clean.csv